

Linear Regression

CMPUT 328

Nilanjan Ray

Computing Science, University of Alberta, Canada

Material source: “Hands-on machine learning with Scikit-Learn and TensorFlow: concepts, tools, and techniques to build intelligent systems,” by Géron, Aurélien.

Linear regression with PyTorch

- We will start with a linear regression “model”
- Next, we need to understand “loss” function for regression task
- Next we will estimate the model by minimizing the loss function
- We will use PyTorch

Quick review: Gradient of a function

Example:


$$f(x, y) = -(\cos^2 x + \cos^2 y)^2$$

$$\nabla f(x, y) = \begin{bmatrix} \frac{\partial f}{\partial x} \\ \frac{\partial f}{\partial y} \end{bmatrix} = \begin{bmatrix} 4(\cos^2(x) + \cos^2(y)) \cos(x) \sin(x) \\ 4(\cos^2(x) + \cos^2(y)) \cos(y) \sin(y) \end{bmatrix}$$

Note 1: f is a function of **two variables**,
so gradient of f is a **two dimensional vector**

Note 2: Gradient (vector) of f points toward the
steepest ascent for f

Note 3: At a (local) minimum of f its gradient
becomes a **zero vector**



Example source: Wikipedia



Quick review: Gradient descent optimization

Start at an initial guess for the optimization variable: \mathbf{x}_0

Iterate until gradient magnitude becomes too small: $\mathbf{x}^{t+1} = \mathbf{x}^t - \alpha \nabla f(\mathbf{x}^t)$

} Gradient descent algorithm

α is called the step-length.



Gradient descent creates a zig-zag path leading to a local minimum of f

Supervised machine learning: the tabular view

Independent variable (aka feature vector)				Prediction / dependent variable
x_1	x_2	x_3	x_4	y
1.2	-3.9	4.0	0	1.6
2.1	2.4	-0.7	-0.2	1.2
...
...
3.2	1.9	0.3
1.4	1.5	?
3.1	2.1	?

Training data:
complete table

Test data:
incomplete table



ML learns to map **x** to **y**

In other words, ML learns
a function, f so that
 $y = f(x)$

The function f is called **prediction function**

Linear prediction: formal setup

Linear prediction function: $y^p = \mathbf{x}\boldsymbol{\theta} + b$ or, $y^p = \sum_{j=1}^m \theta_j x_j + b$

vector equation form

scalar equation form

A training set consists of (\mathbf{x}, y) pairs: $\{(\mathbf{x}_i, y_i)\}_{i=1}^n$

Linear prediction on the training data point i : $y_i^p = \mathbf{x}_i \boldsymbol{\theta} + b$ or, $y_i^p = \sum_{j=1}^m \theta_j x_{i,j} + b$

Loss or cost function (on training data): $L = \frac{1}{2} \sum_{i=1}^n (y_i^p - y_i)^2$

Learning the linear model: find out $\boldsymbol{\theta}$ and b to minimize loss function

Linear regression: A toy example

Let's take a toy example:

x_1	x_2	y
1	2	-1
3	-4	7
6	2	3
-3	5	-4
7	-3	5
4	3	?

This equation $y_i^p = \sum_{j=1}^m \theta_j x_{i,j} + b$

can be written for the toy training set as

We also have responses:

$$\begin{aligned}y_1^p &= \theta_1(1) + \theta_2(2) + b \\y_2^p &= \theta_1(3) + \theta_2(-4) + b \\y_3^p &= \theta_1(6) + \theta_2(2) + b \\y_4^p &= \theta_1(-3) + \theta_2(5) + b \\y_5^p &= \theta_1(7) + \theta_2(-3) + b\end{aligned}$$

$$y_1 = -1, y_2 = 7, y_3 = 3, y_4 = -4, y_5 = 5$$

So, the loss is
$$L = \frac{1}{2} \sum_{i=1}^n (y_i^p - y_i)^2 = \frac{1}{2} [(y_1^p + 1)^2 + (y_2^p - 7)^2 + (y_3^p - 3)^2 + (y_4^p + 4)^2 + (y_5^p - 5)^2]$$

Learning a linear model

For the convenience of math, let us change our linear model a bit:

$$y_i^p = \sum_{j=1}^m \theta_j (x_{i,j} - \bar{x}_j) + b \quad \text{where,} \quad \bar{x}_j = \frac{1}{n} \sum_{i=1}^n x_{i,j}$$

And a slightly modified loss function:

$$L = \underbrace{\frac{1}{2} \sum_{i=1}^n (y_i^p - y_i)^2}_{\text{Data fidelity}} + \underbrace{\frac{\gamma}{2} \sum_{j=1}^m \theta_j^2}_{\text{Regularization}}$$

γ is a **hyper parameter**

Why do we need regularization?

Minimization of linear regression loss function

Regularized loss function:
$$L = \frac{1}{2} \sum_{i=1}^n (y_i^p - y_i)^2 + \frac{\gamma}{2} \sum_{j=1}^m \theta_j^2$$

Taking partial derivative using chain rule:
$$\frac{\partial L}{\partial b} = \sum_{i=1}^n (y_i^p - y_i) \frac{\partial y_i^p}{\partial b} = \sum_{i=1}^n (y_i^p - y_i) \quad \text{because,} \quad \frac{\partial y_i^p}{\partial b} = 1$$

Using $\bar{x}_j = \frac{1}{n} \sum_{i=1}^n x_{i,j}$ and $y_i^p = \sum_{j=1}^m \theta_j (x_{i,j} - \bar{x}_j) + b$ we get:
$$\frac{\partial L}{\partial b} = nb - \sum_{i=1}^n y_i$$

At the minimum of L , $\frac{\partial L}{\partial b} = 0$ So,
$$b = \frac{1}{n} \sum_{i=1}^n y_i = \bar{y}$$

Linear regression: A toy example...continued

Let's take a toy example:

x_1	x_2	y
1	2	-1
3	-4	7
6	2	3
-3	5	-4
7	-3	5
4	3	?

$$b = \frac{1}{n} \sum_{i=1}^n y_i = \bar{y} = \frac{1}{5}(-1 + 7 + 3 - 4 + 5) = 2$$

$$\bar{x}_1 = \frac{1}{n} \sum_{i=1}^n x_{i,1} = \frac{1}{5}(1 + 3 + 6 - 3 + 7) = 2.8$$

$$\bar{x}_2 = \frac{1}{n} \sum_{i=1}^n x_{i,2} = \frac{1}{5}(2 - 4 + 2 + 5 - 3) = 0.4$$

So, using centered data, the prediction equation becomes:

$$y_i^p = \sum_{j=1}^m \theta_j (x_{i,j} - \bar{x}_j) + b = \theta_1 (x_{i,1} - 2.8) + \theta_2 (x_{i,2} - 0.4) + 2$$

So, the loss is

$$\begin{aligned} L &= \frac{1}{2} \sum_{i=1}^n (y_i^p - y_i)^2 = \frac{1}{2} [(y_1^p + 1)^2 + (y_2^p - 7)^2 + (y_3^p - 3)^2 + (y_4^p + 4)^2 + (y_5^p - 5)^2] \\ &= \frac{1}{2} [(\theta_1(1 - 2.8) + \theta_2(2 - 0.4) + 2 + 1)^2 + (\theta_1(3 - 2.8) + \theta_2(-4 - 0.4) + 2 - 7)^2 \\ &\quad + (\theta_1(6 - 2.8) + \theta_2(2 - 0.4) + 2 - 3)^2 + (\theta_1(-3 - 2.8) + \theta_2(5 - 0.4) + 2 + 4)^2 + (\theta_1(7 - 2.8) + \theta_2(-3 - 0.4) + 2 - 5)^2] \end{aligned}$$

Minimization of linear regression loss function...

Regularized loss function:

$$L = \frac{1}{2} \sum_{i=1}^n (y_i^p - y_i)^2 + \frac{\gamma}{2} \sum_{j=1}^m \theta_j^2$$

Taking partial derivative of L using chain rule:

$$\frac{\partial L}{\partial \theta_j} = \sum_{i=1}^n (y_i^p - y_i) \frac{\partial y_i^p}{\partial \theta_j} + \gamma \theta_j$$

Using $y_i^p = \sum_{k=1}^m \theta_k (x_{i,k} - \bar{x}_k) + b$, $b = \bar{y}$ and $\frac{\partial y_i^p}{\partial \theta_j} = x_{i,j} - \bar{x}_j$

We get:
$$\frac{\partial L}{\partial \theta_j} = \sum_{i=1}^n \left(\sum_{k=1}^m \theta_k (x_{i,k} - \bar{x}_k) + \bar{y} - y_i \right) (x_{i,j} - \bar{x}_j) + \gamma \theta_j$$

Linear regression: A toy example...continued

Let's take a toy example:

x_1	x_2	y
1	2	-1
3	-4	7
6	2	3
-3	5	-4
7	-3	5
4	3	?

Note: For this problem I did not assume any regularization


$$\frac{\partial L}{\partial \theta_j} = \sum_{i=1}^n \left(\sum_{k=1}^m \theta_k (x_{i,k} - \bar{x}_k) + \bar{y} - y_i \right) (x_{i,j} - \bar{x}_j) + \gamma \theta_j$$

$$\begin{aligned} \frac{\partial L}{\partial \theta_1} &= (\theta_1(1 - 2.8) + \theta_2(2 - 0.4) + 2 + 1)(1 - 2.8) \\ &+ (\theta_1(3 - 2.8) + \theta_2(-4 - 0.4) + 2 - 7)(3 - 2.8) \\ &+ (\theta_1(6 - 2.8) + \theta_2(2 - 0.4) + 2 - 3)(6 - 2.8) \\ &+ (\theta_1(-3 - 2.8) + \theta_2(5 - 0.4) + 2 + 4)(-3 - 2.8) \\ &+ (\theta_1(7 - 2.8) + \theta_2(-3 - 0.4) + 2 - 5)(7 - 2.8) \end{aligned}$$

$$\begin{aligned} \frac{\partial L}{\partial \theta_2} &= (\theta_1(1 - 2.8) + \theta_2(2 - 0.4) + 2 + 1)(2 - 0.4) \\ &+ (\theta_1(3 - 2.8) + \theta_2(-4 - 0.4) + 2 - 7)(-4 - 0.4) \\ &+ (\theta_1(6 - 2.8) + \theta_2(2 - 0.4) + 2 - 3)(2 - 0.4) \\ &+ (\theta_1(-3 - 2.8) + \theta_2(5 - 0.4) + 2 + 4)(5 - 0.4) \\ &+ (\theta_1(7 - 2.8) + \theta_2(-3 - 0.4) + 2 - 5)(-3 - 0.4) \end{aligned}$$

Minimization of linear regression loss function...

$$\frac{\partial L}{\partial \theta_j} = \sum_{i=1}^n \left(\sum_{k=1}^m \theta_k (x_{i,k} - \bar{x}_k) + \bar{y} - y_i \right) (x_{i,j} - \bar{x}_j) + \gamma \theta_j$$

simplification


Gradient of L :

$$\nabla L = \left[\sum_{i=1}^n (\mathbf{x}_i - \bar{\mathbf{x}})^T (\mathbf{x}_i - \bar{\mathbf{x}}) \right] \boldsymbol{\theta} + \gamma \boldsymbol{\theta} - \sum_{i=1}^n (y_i - \bar{y}) (\mathbf{x}_i - \bar{\mathbf{x}})^T$$

where $\mathbf{x}_i = [x_{i,1} \quad \dots \quad x_{i,m}]$, $\bar{\mathbf{x}} = [\bar{x}_1 \quad \dots \quad \bar{x}_m]$ and $\boldsymbol{\theta} = [\theta_1 \quad \dots \quad \theta_m]^T$

More simplified form:

$$\nabla L = (X^T X + \gamma I) \boldsymbol{\theta} - X^T \mathbf{y}$$

where matrix X is defined as: $X = \begin{bmatrix} \mathbf{x}_1 - \bar{\mathbf{x}} \\ \vdots \\ \mathbf{x}_n - \bar{\mathbf{x}} \end{bmatrix}$ and vector \mathbf{y} is defined as: $\mathbf{y} = \begin{bmatrix} y_1 - \bar{y} \\ \vdots \\ y_n - \bar{y} \end{bmatrix}$

and I is an identity matrix of size m -by- m

Equating gradient of L to zero vector and solving gives us:

$$\boldsymbol{\theta} = (X^T X + \gamma I)^{-1} X^T \mathbf{y}$$

Linear regression: A toy example...finally!

Let's take a toy example:

x_1	x_2	y
1	2	-1
3	-4	7
6	2	3
-3	5	-4
7	-3	5
4	3	?

$$X = \begin{bmatrix} \mathbf{x}_1 - \bar{\mathbf{x}} \\ \vdots \\ \mathbf{x}_n - \bar{\mathbf{x}} \end{bmatrix} = \begin{bmatrix} 1 - 2.8 & 2 - 0.4 \\ 3 - 2.8 & -4 - 0.4 \\ 6 - 2.8 & 2 - 0.4 \\ -3 - 2.8 & 5 - 0.4 \\ 7 - 2.8 & -3 - 0.4 \end{bmatrix} \quad \mathbf{y} = \begin{bmatrix} y_1 - \bar{y} \\ \vdots \\ y_n - \bar{y} \end{bmatrix} = \begin{bmatrix} -1 - 2 \\ 7 - 2 \\ 3 - 2 \\ -4 - 2 \\ 5 - 2 \end{bmatrix}$$

$$\boldsymbol{\theta} = (X^T X)^{-1} X^T \mathbf{y} = \begin{bmatrix} 0.3580 \\ -0.8535 \end{bmatrix}$$

So, finally the prediction for the test data point

$$? = \sum_{j=1}^m \theta_j (x_j - \bar{x}_j) + b = 0.3580(4 - 2.8) - 0.8535(3 - 0.4) + 2 = 0.2105$$

Note: For this problem I did not assume any regularization

Linear regression by gradient descent

$$\boldsymbol{\theta} = (X^T X + \gamma I)^{-1} X^T \mathbf{y}$$

If the data does not fit into the memory, you cannot compute $\boldsymbol{\theta}$ directly with the above formula; you apply **gradient descent** to compute it (approximately).

$$\nabla L(\boldsymbol{\theta}; X, \mathbf{y}) = X^T X \boldsymbol{\theta} + \gamma \boldsymbol{\theta} - X^T \mathbf{y}$$

Guess a starting value for $\boldsymbol{\theta} = \boldsymbol{\theta}_0$

Initialize learning rate and regularization parameter: α, γ

Iterate for $t = 0, 1, \dots$

Consider a subset of data (X_t, \mathbf{y}_t)

Update: $\boldsymbol{\theta}_{t+1} = \boldsymbol{\theta}_t - \alpha \nabla L(\boldsymbol{\theta}_t; X_t, \mathbf{y}_t)$

Gradient descent algorithm

Also known as **batch update** or **batch method**

Derivation using vector calculus

Regularized loss function: $L = \frac{1}{2} \sum_{i=1}^n (y_i^p - y_i)^2 + \frac{\gamma}{2} \sum_{j=1}^m \theta_j^2$ or, $L = \frac{1}{2} \sum_{i=1}^n (y_i^p - y_i)^2 + \frac{\gamma}{2} \boldsymbol{\theta}^T \boldsymbol{\theta}$

Using vector calculus: $\nabla L = \sum_{i=1}^n (y_i^p - y_i) \nabla y_i^p + \frac{\gamma}{2} \nabla (\boldsymbol{\theta}^T \boldsymbol{\theta})$

Use vector differentiation to: $y_i^p = (\mathbf{x}_i - \bar{\mathbf{x}}) \boldsymbol{\theta} + \bar{y}$ and get: $\nabla y_i^p = (\mathbf{x}_i - \bar{\mathbf{x}})^T$

Also note, using vector differentiation rule: $\nabla (\boldsymbol{\theta}^T \boldsymbol{\theta}) = 2\boldsymbol{\theta}$

$$\nabla L = \underbrace{\begin{bmatrix} \mathbf{x}_1 - \bar{\mathbf{x}} \\ \vdots \\ \mathbf{x}_n - \bar{\mathbf{x}} \end{bmatrix}^T}_{\text{"centered data"}} \underbrace{\begin{bmatrix} y_1^p - y_1 \\ \vdots \\ y_n^p - y_n \end{bmatrix}}_{\text{"error"}} + \gamma \boldsymbol{\theta}$$

MNIST Dataset



Classify images into digits

Each image is **28x28**

10 labels

55,000 training images

5,000 validation images

10,000 test images.

Linear regression on MNIST dataset



Small 28 pixels-by-28 pixels images of hand written digits

The visual recognition problem definition:
to recognize the digit from an image

Our very first line of attack would be to
use linear regression.

Feature dimension, $m = 28 * 28 = 784$

Let's look at our PyTorch implementation

Pixel values (feature)				Digit
x_1	x_2	...	x_{784}	y
0.1	0.3	...	0.0	0
0.2	0.1	...	0.5	1
...
...
0.0	0.98	...	0.8	9
0.5	0.25	...	0.36	?
0.1	0.95	...	0.1	?